

```
library(stats) library(caret) library(glmnet)

df <- read.csv("tecator.csv") n <- dim(df)[1] source("split_data.R") sets <-
split_data(data_frame = df, k = 2, seed = 12345, distributions = c(1,1)) train <- sets[[1]] test
<- sets[[2]]

#Channels as features train_features <- train[1:(ncol(train)-3)] test_features <-
test[1:(ncol(test)-3)]

#fitting linear regression model lin_model <- lm(train$Fat ~. +0, data = train_features ) #
+0 takes away the intercept

#predictions on the training data and test data pred_train <- predict(lin_model,
train_features) pred_test <- predict(lin_model, test_features)
```

Training and test errors

```
linear_mse_train <- mean((trainFat - pred_train)^2) linear_mse_test <- mean((testFat -
pred_test) ^ 2)

summary(lin_model)

#=====part1=====
=====

#Lasso regression X_train <- as.matrix(train_features) X_test <- as.matrix(test_features)
```

Fit LASSO regression model on the training data

```
lasso_model_train <- cv.glmnet(X_train, train$Fat , alpha = 1) # Setting alpha = 1 for LASSO
```

LASSO Cost Function

```
lasso_cost_function <- function(X,y, beta, lambda) { N <- length(y) prediction <- X %% beta
residuals <- y - prediction mse_term <- sum(residuals^2) / (2 * N) l1_penalty <- lambda *
sum(abs(beta)) cost <- mse_term + l1_penalty return(cost) }
```

Example usage of the cost function

```
optimal_lambda <- as.matrix(lasso_model_train$lambda.min) optimal_beta <-
coef(lasso_model_train, s = optimal_lambda) #The discrepancy in the dimensions suggests
that the optimal_beta vector includes an additional element, #which typically corresponds
to the intercept term. #In glmnet models, the intercept is included by default.
```

```
coefficients_optimal <- as.matrix(tail(optimal_beta, -1) ) # Exclude the first element (intercept)
```

Ensure that X_train_lasso and optimal_beta have compatible dimensions

```
if (ncol(X_train) == length(coefficients_optimal)) { cost <- lasso_cost_function(X_train, train$Fat, coefficients_optimal, optimal_lambda) cat("LASSO Cost Function Value:", cost, "") } else { cat("Dimensions of X_train_lasso and optimal_beta are not compatible for matrix multiplication.") }
```

```
#=====part2=====
```

Fit LASSO regression model on the training data

```
summary(lasso_model_train)
```

Plot the coefficient paths

```
plot(lasso_model_train$glmnet.fit, "lambda", main = "LASSO Coefficient Paths", xlab = "Log(lambda)", ylab = "Coefficients")
```

Identify the value of lambda for a model with only three features

```
desired_num_features <- 3 lambda_index <- which.min(lasso_model_train$glmnet.fit$dev.ratio > desired_num_features / ncol(X_train))
```

Add a vertical line at the selected lambda value

```
abline(v = log(lasso_model_train$glmnet.fit$lambda[lambda_index]), col = "red", lty = 2)
```

```
#=====part3=====
```

Fit Ridge regression model

```
ridge_model_train <- cv.glmnet(X_train, train$Fat, alpha = 0) # Setting alpha = 0 for ridge
```

Plotting regression coefficients vs. $\log(\lambda)$

```
plot(ridge_model_train$glmnet.fit, "lambda", main = "Ridge Coefficient Paths", xlab =  
"Log(lambda)", ylab = "Coefficients")
```

```
#=====part4=====
```

```
#cross-validation to compute optimal lasso model cv_lasso_model <- cv.glmnet(X_test,  
test$Fat, alpha = 1)
```

```
plot(cv_lasso_model)
```

Optimal λ and number of variables

```
opt_lambda <- cv_lasso_model$lambda.min  
n_variables_chosen <-  
sum(coef(cv_lasso_model, s = optimal_lambda) != 0)  
lambda_four <-  
cv_lasso_model$cvm[which.min(abs(log(cv_lasso_model$lambda) -  
(-4)))]  
lambda_optimal <-  
cv_lasso_model$cvm[which.min(abs(log(cv_lasso_model$lambda) - opt_lambda)]]
```

```
test_pred_lasso <- predict(cv_lasso_model, newx = X_test, s = opt_lambda)  
plot(test$Fat, test_pred_lasso, main = "Actual vs. Predicted (LASSO)", xlab = "Actual", ylab = "Predicted",  
col = c("blue", "red"))  
legend("topleft", legend = c("Actual", "Predicted"), col = c("blue",  
"red"), pch = 1)
```

```
#=====part5=====
```